

D* Lite Pathfinding

Gemmin Sugiura

April 2025

1 Introduction

This document provides a (hopefully simpler) explanation of the D* Lite pathfinding algorithm. It is based on the original research paper D* Lite by Sven Koenig and Maxim Likhachev. Readers interested in a more in-depth understanding are encouraged to refer to their published work. D* Lite pathfinding addresses the problem with dynamic graphs (i.e. when vertices' traversability changes). It is the best algorithm if the target position does not change often.

2 Basic Definitions

- $G(V, E)$ denotes an undirected graph, where V denotes the set of vertices and E denotes the set of edges.
- $\text{NEIGH}(v) \subseteq V$ denotes the set of vertices that can be reached directly from v .
- $v_{\text{start}} \in V$ denotes the start vertex (i.e. the agent's current position). It is dynamic since the agent moves.
- $v_{\text{end}} \in V$ is the target vertex. It is static.
- $c(u, v) : E \rightarrow \mathbb{R}^+$ denotes the edge cost from vertex u to vertex v . G is undirected implying $c(u, v) = c(v, u)$.
- $u \in V$ is traversable iff $c(u, v) < \infty, \forall v \in \text{NEIGH}(u)$ (like toggle).
- $u \in V$ is intraversable iff $c(u, v) = \infty, \forall v \in \text{NEIGH}(u)$.
- $g(v) : V \rightarrow \mathbb{R}^+$ denotes the estimate length of the shortest path from v_{start} to v .
- $g^*(v) : V \rightarrow \mathbb{R}^+$ denotes the actual length of the shortest path from v_{start} to v .
- $h(v) : V \rightarrow \mathbb{R}^+$ is the heuristic function approximating the distance from v to v_{goal} . $h(v)$ is admissible and follows the triangle inequality.
- H denotes the heap (I use priority queue synonymously).

3 Lifelong Planning A* (LPA*) Pathfinding

Understanding LPA* is a prerequisite for learning D* Lite, as the latter builds upon the former's key concepts. There are three main ideas:

1. Right hand Side (RHS)
2. Local Consistency
3. Heap and Keys

3.1 Right Hand Side (RHS)

1. The $rhs(v)$ provides an alternative estimate of $g(v)$, computed using the vertex's local neighborhood information. This value serves as a "reality check." While $g(v)$ stores the current estimated cost from the start, $rhs(v)$ asks: Based on my neighbors' actual costs ($g(u)$) and their connection to me ($c(u, v)$), what should my cost reasonably be? Typically $g(v) \geq rhs(v)$, as the one-step look-ahead through neighbors often reveals better paths than the current $g(v)$ estimate. Formally:

$$rhs(v) = \begin{cases} \min_{u \in \text{NEIGH}(v)} (g(u) + c(u, v)), & \text{if } v \neq v_{start} \\ 0 & \text{else} \end{cases}$$

3.2 Local Consistency

A vertex v is locally consistent when its stored cost $g(v)$ matches the best estimate of its neighbors: $g(v) = rhs(v)$. When all vertices are consistent, the $g = g^*$. Therefore, the goal is to make all vertices consistent.

| Type | Condition | Meaning |
|-----------------|-----------------|--|
| Overconsistent | $g(v) > rhs(v)$ | Current cost estimate is too high (too pessimistic). |
| Underconsistent | $g(v) < rhs(v)$ | Current cost is invalid. |

3.3 Min Heap and Keys

Only locally inconsistent vertices need processing. Thus, the heap should consist of vertices that are locally inconsistent, that is, $g(v) \neq rhs(v)$. Each vertex in the heap should have a priority, which we will call the key. Each key is a pair of two values $[k_1, k_2] = [\min(g(v), rhs(v)) + h(v, v_{goal}), \min(g(v), rhs(v))]$. The best way to understand the key is to first understand that $\min(g(v), rhs(v))$ is just an estimate of the shortest path length from the start to some vertex. Then $k_1 = g(v) + h(v)$ is equal to the key used in A*, and k_2 is simply a tiebreaker. The priority queue must maintain: $k(v) = \text{calculateKey}(v), \forall v \in H$, where $k(v)$ is the stored key, and $\text{calculateKey}(v)$ is the actual priority.

3.4 Pseudocode

- 1: **procedure** CALCULATEKEY(v)
- 2: **return** $[\min(g(v), rhs(v)) + h(v, v_{goal}); \min(g(v), rhs(v))]$

```

3: procedure INITIALIZE
4:    $H = \emptyset$ 
5:   for all  $v \in V$  do                                     ▷ In practice, use map
6:      $rhs(v) = g(v) = \infty$ 
7:      $rhs(v_{start}) = 0$ 
8:      $H.Insert(v_{start}, CalculateKey(s_{start}))$ 
9: procedure UPDATEVERTEX( $v$ )
10:  if  $v \neq v_{start}$  then
11:     $rhs(v) = \min_{v' \in NEIGH(v)} (g(v') + c(v', v))$ 
12:  if  $v \in H$  then
13:     $H.Remove(v)$ 
14:  if  $g(v) \neq rhs(v)$  then
15:     $H.Insert(v, CalculateKey(v))$ 
16: procedure COMPUTESHORTESTPATH
17:  while  $H.TopKey() < CalculateKey(s_{goal})$  OR  $rhs(s_{goal}) \neq g(s_{goal})$  do
18:     $v = H.Pop()$ 
19:    if  $g(v) > rhs(v)$  then
20:       $g(v) = rhs(v)$ 
21:      for all  $s \in NEIGH(v)$  do
22:        UpdateVertex( $s$ )
23:    else
24:       $g(v) = \infty$ 
25:      for all  $s \in NEIGH(v) \cup \{v\}$  do
26:        UpdateVertex( $s$ )
27: procedure MAIN
28:  Initialize()
29:  loop
30:    ComputeShortestPath()
31:    Wait for changes in edge costs
32:    for all directed edges  $(u, v)$  with changed edge costs do
33:      Update the edge cost  $c(u, v)$ 
34:      UpdateVertex( $v$ )

```

4 D* Lite Pathfinding

Unlike LPA*, D* Lite performs a backward search from v_{goal} , making $g(v)$ represent the estimated shortest path cost from v to v_{goal} . Like LPA*, I divided the learning into three sections.

1. Key Mechanics
2. Error Analysis
3. Correction Method

4.1 Key Mechanics

The key in D* Lite is slightly different: $k_1(v) = \min(g(v), rhs(v)) + h(v_{\text{start}}, v) + k_m$, where v_{start} is the agent's position. $\min(g(v), rhs(v))$ remains a constant, but the heuristic $h(v_{\text{start}}, v)$ changes with respect to the agent's position. Consequently, the keys stored in the priority queue become obsolete because the stored keys were calculated using the old agent's position. The naive fix is to recalculate the priorities of all vertices every time the agent moves.

4.2 Error Analysis

Intuitively, as the agent moves along or near the optimal path toward the target vertex, $h(v_{\text{start}}, v)$ should decrease, implying that the stored keys become overestimates of the true priorities. Let $\varepsilon = \Delta k_1$ denote this error. Since the stored keys were calculated using $h(v_{\text{old}}, v)$, and the actual priorities use $h(v_{\text{curr}}, v)$, $\varepsilon = \Delta k_1 = \Delta h = h(v_{\text{old}}, v) - h(v_{\text{curr}}, v)$. ε can be upper-bounded by using the triangle inequality.

$$\begin{aligned} h(v_{\text{old}}, v) &\leq h(v_{\text{old}}, v_{\text{curr}}) + h(v_{\text{curr}}, v) && \text{(Triangle inequality)} \\ \implies h(v_{\text{old}}, v) - h(v_{\text{curr}}, v) &\leq h(v_{\text{old}}, v_{\text{curr}}) && \text{(Rearranged form)} \end{aligned}$$

4.3 Correction Method

Suppose that an agent moves from v to v' following an optimal path and detects changes in the edge cost(s), forcing new priorities to be computed. Then, from the result above, the new priorities will have decreased by at most $h(v_{\text{old}}, v_{\text{curr}})$. In other words, when edge cost(s) change, the new priorities are $h(v, v')$ too small relative to $k_1(v)$, $\forall v \in H$. Thus, we could add $k_m = h(v, v')$ to the new priorities.

4.4 Pseudocode

```

1: procedure CALCULATEKEY( $v$ )
2:   return  $[\min(g(v), rhs(v)) + h(v, v_{\text{start}}) + k_m, \min(g(v), rhs(v))]$ 
3: procedure UPDATEVERTEX( $v$ )
4:   if  $v \in H$  and  $g(v) \neq rhs(v)$  then
5:      $H.\text{update}(v, \text{calculateKey}(v))$ 
6:   else if  $v \notin H$  and  $g(v) \neq rhs(v)$  then
7:      $H.\text{insert}(v, \text{calculateKey}(v))$ 
8:   else if  $v \in H$  and  $g(v) = rhs(v)$  then
9:      $H.\text{remove}(v)$ 
10: procedure COMPUTESHORTESTPATH
11:   while  $H.\text{topKey}() < \text{calculateKey}(v_{\text{start}})$  or  $rhs(v_{\text{start}}) > g(v_{\text{start}})$  do
12:      $\{v, k_{\text{old}}\} \leftarrow H.\text{top}(), H.\text{topKey}()$ 
13:      $k_{\text{new}} \leftarrow \text{calculateKey}(v)$ 
14:     if  $k_{\text{old}} < k_{\text{new}}$  then ▷ Underestimate
15:        $H.\text{update}(v, k_{\text{new}})$ 
16:     else if  $g(v) > rhs(v)$  then ▷ Locally Overconsistent
17:        $g(v) \leftarrow rhs(v)$ 

```

```

18:          $H.remove(v)$ 
19:         for all  $u \in \text{NEIGH}(v)$  do
20:             if  $u \neq v_{\text{goal}}$  then
21:                  $rhs(u) \leftarrow \min(rhs(u), c(u, v) + g(v))$ 
22:                  $updateVertex(u)$ 
23:         else ▷ Locally Underconsistent
24:              $oldG \leftarrow g(v)$ 
25:              $g(v) \leftarrow \infty$ 
26:             for all  $u \in \text{pred}(v) \cup \{v\}$  do
27:                 if  $rhs(u) = c(u, v) + oldG$  then
28:                      $rhs(u) \leftarrow computeRhs(u)$ 
29:              $updateVertex(u)$ 
30: procedure MAIN
31:      $v_{\text{last}} \leftarrow v_{\text{start}}$ 
32:     initialize()
33:     computeShortestPath()
34:     while  $v_{\text{start}} \neq v_{\text{goal}}$  do
35:          $v_{\text{start}} \leftarrow \arg \min_{v' \in \text{NEIGH}(v_{\text{last}})} (c(v_{\text{last}}, v') + g(v'))$ 
36:         Move to  $v_{\text{start}}$ 
37:         if edge costs changed then
38:              $k_m \leftarrow k_m + h(v_{\text{last}}, v_{\text{start}})$ 
39:              $v_{\text{last}} \leftarrow v_{\text{start}}$ 
40:             for all  $(u, v)$  with changed edge costs do
41:                  $c_{\text{old}} \leftarrow c(u, v)$ 
42:                 Update  $c(u, v)$ 
43:                 if  $c_{\text{old}} > c(u, v)$  then
44:                     if  $u \neq v_{\text{goal}}$  then
45:                          $rhs(u) \leftarrow \min(rhs(u), c(u, v) + g(v))$ 
46:                     else if  $rhs(u) = c_{\text{old}} + g(v)$  then
47:                         if  $u \neq v_{\text{goal}}$  then
48:                              $rhs(u) \leftarrow \min_{u' \in \text{NEIGH}(u)} (c(u, u') + g(u'))$ 
49:                  $updateVertex(u)$ 
50:             computeShortestPath()

```

4.4.1 Explanation of computeShortestPath()

When a vertex v is popped from the priority queue, there are three cases:

1. $k(v)$ is an underestimate of the true priority. Therefore, the vertex should be reinserted into the priority queue with its corrected key (line 15).
2. v is locally overconsistent meaning $g(v) > rhs(v)$. Simply, update $g(v)$ with $rhs(v)$ to make v locally consistent (lines 16). Since changing $g(v)$ can affect the rhs values of its neighbors, fix the neighbors' rhs values to reflect the updated $g(v)$ (lines 20-21).
3. v is locally underconsistent meaning $g(v) < rhs(v)$ (lines 23). Case 3 will only happen when v becomes intraversable. Since $g(v) < rhs(v)$, $g(v)$ is too optimistic. Therefore, $g(v)$ is reset to ensure the recomputation of

$g(v)$, that is, to make v locally overconsistent (line 25). Since $g_{\text{old}}(v)$ is outdated, the rhs values of its neighbors may be wrong. Thus, for each of the neighboring vertices, we ask if $g_{\text{old}}(v)$ was used to calculate rhs (line 27). If so, fix rhs using its definition (lines 27-28).

4.4.2 Explanation of `main()`

While the agent is not at v_{goal} , the agent moves greedily toward v_{goal} (line 35). If traversability changes, we use the heap ordering trick to account for new priorities that must be computed (lines 38-39). Some edge costs can change, indicating changes in traversability. Therefore, when some edge costs change, the rhs values should be updated to reflect these changes. Changes in the rhs values mean that new priorities must be calculated. Thus, the offset value k_m should also be updated before calling `updateVertex(v)` (line 49). For each changed edge costs, there are two cases. If the edge cost decreased,